Basic Spin Manual

Contents

Spin is a tool for analyzing the logical consistency of concurrent systems, specifically of data communication protocols. The system is described in a modeling language called `Promela` (Process Meta Language). The language allows for the dynamic creation of concurrent processes. Communication via message channels can be defined to be synchronous (i.e., rendezvous), or asynchronous (i.e., buffered).

Given a model system specified in `Promela`, *Spin* can either perform random simulations of the system's execution or it can generate a C program that performs an efficient online verification of the system's correctness properties. During simulation and verification *Spin* checks for the absence of deadlocks, unspecified receptions, and unexecutable code. The verifier can also be used to verify the correctness of system invariants, it can find non-progress execution cycles, and it can verify correctness properties expressed in next-time free linear temporal logic formulae.

The verifier is setup to be efficient and to use a minimal amount of memory. An exhaustive verification performed by *Spin* can establish with mathematical certainty whether or not a given behavior is error-free. Very large verification problems, that can ordinarily not be solved within the constraints of a given computer system, can be attacked with a frugal ``bit state storage'' technique, also known as *supertrace*. With this method the state space can be collapsed to a small number of bits per reachable system state, with minimal side-effects.

The first part of this memo gives an introduction to `Promela`, the second part discusses the usage of *Spin*. A concise reference to `Promela` can be found in a [companion document](#). In the Appendix some examples are used to illustrate the construction of a basic `Promela` model for *Spin* verification.

This manual discusses only the most basic use of *Spin*. It does not discuss extensions to the language that are part of more recent versions of the tool (see the notes at the end of this manual). It also does not discuss the builtin support for the verification of linear temporal logic formulae.

Modeling Language

`Promela` is a verification modeling language. It provides a vehicle for making abstractions of protocols (or distributed systems in general) that suppress details that are unrelated to process interaction. The intended use of *Spin* is to verify fractions of process behavior, that for one reason or another are considered suspect. The relevant behavior is modeled in `Promela` and verified. A complete verification is therefore typically performed in a series of steps, with the construction of increasingly detailed `Promela` models. Each model can be verified with *Spin* under different types of assumptions about the environment (e.g., message loss, message duplications etc). Once the correctness of a model has been established with *Spin*, that fact can be used in the construction and verification of all subsequent models.

`Promela` programs consist of *processes*, message *channels*, and *variables*. Processes are global objects. Message channels and variables can be declared either globally or locally within a process. Processes specify behavior, channels and global variables define the environment in which the processes run.

Executability

In `Promela` there is no difference between conditions and statements, even isolated boolean conditions can be used as statements. The execution of every statement is conditional on its *executability* . Statements are either executable or blocked. The executability is the basic means of synchronization. A process can wait for an event to happen by waiting for a statement to become executable. For instance, instead of writing a busy wait loop:

```
while (a != b)
        skip    /* wait for a==b */
```
one can achieve the same effect in `Promela` with the statement
```
(a == b)
```
A condition can only be executed (passed) when it holds. If the condition does not hold, execution blocks until it does.

Variables are used to store either global information about the system as a whole, or information local to one specific process, depending on where the declaration for the variable is placed. The declarations

```
bool flag;
int state;
byte msg;
```
define variables that can store integer values in three different ranges. The scope of a variable is global if it is declared outside all process declarations, and local if it is declared within a process declaration.

Data Types

The table below summarizes the basic data types, sizes, and typical value ranges on a 32-bit wordsize computer. (See also datatypes.)

Table 1 - Data Types

Typename	C-equivalent	Macro in limits.h	Typical Range
bit or bool	-	-	0..1
byte	bit-field	CHAR_BIT (width in bits)	0..255
short	uchar	SHRT_MIN..SHRT_MAX	-2^15 - 1 .. 2^15 - 1
int	short	INT_MIN..INT_MAX	-2^31 - 1 .. 2^31 - 1
	int		

The names `bit` and `bool` are synonyms for a single bit of information. A `byte` is an unsigned quantity that can store a value between 0 and 255. `short`s and `int`s are signed quantities that differ only in the range of values they can hold.

An `mtype` variable can be assigned symbolic values that are declared in an `mtype = { ... }` statement, to be discussed below.

Array Variables

Variables can be declared as arrays. For instance,

```
byte state[N]
```

declares an array of `N` bytes that can be accessed in statements such as

```
state[0] = state[3] + 5 * state[3*2/n]
```

where `n` is a constant or a variable declared elsewhere. The index to an array can be any expression that determines a unique integer value. The effect of an index value outside the range 0.. N-1 is undefined; most likely it will cause a runtime error. (Multi-dimensional arrays can be defined indirectly with the help of the `typedef` construct, see WhatsNew.html, Section 2.1.7)

So far we have seen examples of variable declarations and of two types of statements: boolean conditions and assignments. Declarations and assignments are always *executable*. Conditions are only executable when they hold.

Process Types

The state of a variable or of a message channel can only be changed or inspected by processes. The behavior of a process is defined in a `proctype` declaration. The following, for instance, declares a process with one local variable `state`.

```
proctype A()
{       byte state;

        state = 3
}
```

The process type is named `A`. The body of the declaration is enclosed in curly braces. The declaration body consists of a list of zero or more declarations of local variables and/or statements. The declaration above contains one local variable declaration and a single statement: an assignment of the value 3 to variable `state`.

The semicolon is a statement *separator* (not a statement terminator, hence there is no semicolon after the last statement). `Promela` accepts two different statement separators:

an arrow `` `->` ``and the semicolon `` `;` ``. The two statement separators are equivalent. The arrow is sometimes used as an informal way to indicate a causal relation between two statements. Consider the following example.

```
byte state = 2;

proctype A()
{       (state == 1) -> state = 3
}

proctype B()
{       state = state - 1
}
```

In this example we declared two types of processes, A and B. Variable `state` is now a global, initialized to the value two. Process type A contains two statements, separated by an arrow. In the example, process declaration B contains a single statement that decrements the value of the state variable by one. Since the assignment is always executable, processes of type B can always complete without delay. Processes of type A, however, are delayed at the condition until the variable `state` contains the proper value.

Process Instantiation

A `proctype` definition only declares process behavior, it does not execute it. Initially, in the `Promela` model, just one process will be executed: a process of type `init`, that must be declared explicitly in every `Promela` specification. The smallest possible `Promela` specification, therefore, is:

```
init { skip }
```

where `skip` is a dummy, null statement. More interestingly, however, the initial process can initialize global variables, and instantiate processes. An `init` declaration for the above system, for instance, could look as follows.

```
init
{       run A(); run B()
}
```

`run` is used as a unary operator that takes the name of a process type (e.g. A). It is executable only if a process of the type specified can be instantiated. It is unexecutable if this cannot be done, for instance if too many processes are already running.

The run statement can pass parameter values of all basic data types to the new process. The declarations are then written, for instance, as follows:

```
proctype A(byte state; short foo)
{
        (state == 1) -> state = foo
}

init
{
        run A(1, 3)
}
```

Data arrays or process types can not be passed as parameters. As we will see below, there is just one other data type that can be used as a parameter: a message channel.

Run statements can be used in any process to spawn new processes, not just in the initial process. Processes are created with the `run` statements. An executing process disappears again when it terminates (i.e., reaches the end of the body of its process type declaration), but not before all processes that it started have terminated.

With the `run` statement we can create any number of copies of the process types `A` and `B`. If, however, more than one concurrent process is allowed to both read and write the value of a global variable a well-known set of problems can result; for example see [2]. Consider, for instance, the following system of two processes, sharing access to the global variable `state`.

```
byte state = 1;

proctype A()
{       byte tmp;
        (state==1) -> tmp = state; tmp = tmp+1; state = tmp
}

proctype B()
{       byte tmp;
        (state==1) -> tmp = state; tmp = tmp-1; state = tmp
}

init
{       run A(); run B()
}
```

If one of the two processes completes before its competitor has started, the other process will block forever on the initial condition. If both pass the condition simultaneously, both will complete, but the resulting value of `state` is unpredictable. It can be any of the values 0, 1, or 2.

Many solutions to this problem have been considered, ranging from an abolishment of global variables to the provision of special machine instructions that can guarantee an indivisible test and set sequence on a shared variable. The example below was one of the first solutions published. It is due to the Dutch mathematician Dekker. It grants two processes mutually exclusion access to an arbitrary *critical section* in their code, by manipulation three additional global variables. The first four lines in the `Promela` specification below are C-style macro definitions. The first two macros define `true` to be a constant value equal to 1 and *false* to be a constant 0. Similarly, *Aturn* and *Bturn* are defined as constants.

```
#define true    1
#define false   0
#define Aturn   false
#define Bturn   true

bool x, y, t;

proctype A()
{       x = true;
        t = Bturn;
        (y == false || t == Aturn);
        /* critical section */
        x = false
}

proctype B()
```

```
{        y = true;
         t = Aturn;
         (x == false || t == Bturn);
         /* critical section */
         y = false
}

init
{        run A(); run B()
}
```

The algorithm can be executed repeatedly and is independent of the relative speeds of the two processes.

Atomic Sequences

In `Promela` there is also another way to avoid the *test and set* problem: `atomic` sequences. By prefixing a sequence of statements enclosed in curly braces with the keyword `atomic` the user can indicate that the sequence is to be executed as one indivisible unit, non-interleaved with any other processes. It causes a run-time error if any statement, other than the first statement, blocks in an atomic sequence. This is how we can use atomic sequences to protect the concurrent access to the global variable `state` in the earlier example.

```
byte state = 1;

proctype A()
{        atomic {
            (state==1) -> state = state+1
         }
}

proctype B()
{        atomic {
            (state==1) -> state = state-1
         }
}

init
{        run A(); run B()
}
```

In this case the final value of `state` is either zero or two, depending on which process executes. The other process will be blocked forever.

Atomic sequences can be an important tool in reducing the complexity of verification models. Note that atomic sequence restricts the amount of interleaving that is allowed in a distributed system. Otherwise untractable models can be made tractable by, for instance, labeling all manipulations of local variables with atomic sequences. The reduction in complexity can be dramatic.

Message Passing

Message channels are used to model the transfer of data from one process to another. They are declared either locally or globally, for instance as follows:

```
chan qname = [16] of { short }
```

This declares a channel that can store up to 16 messages of type `short`. Channel names can be passed from one process to another via channels or as parameters in process instantiations. If the messages to be passed by the channel have more than one field, the declaration may look as follows:

```
chan qname = [16] of { byte, int, chan, byte }
```

This time the channel stores up to sixteen messages, each consisting of two 8-bit values, one 32-bit value, and a channel name.

The statement

```
qname!expr
```

sends the value of expression `expr` to the channel that we just created, that is: it appends the value to the tail of the channel.

```
qname?msg
```

receives the message, it retrieves it from the head of the channel, and stores it in a variable `msg`. The channels pass messages in first-in-first-out order. In the above cases only a single value is passed through the channel. If more than one value is to be transferred per message, they are specified in a comma separated list

```
qname!expr1,expr2,expr3
qname?var1,var2,var3
```

It is an error to send or receive either more or fewer parameters per message than was declared for the message channel used.

By convention, the first message field is often used to specify the message type (i.e. a constant). An alternative, and equivalent, notation for the send and receive operations is therefore to specify the message type, followed by a list of message fields enclosed in braces. In general:

```
qname!expr1(expr2,expr3)
qname?var1(var2,var3)
```

The send operation is executable only when the channel addressed is not full. The receive operation, similarly, is only executable when the channel is non empty. Optionally, some of the arguments of the receive operation can be constants:

```
qname?cons1,var2,cons2
```

in this case, a further condition on the executability of the receive operation is that the value of all message fields that are specified as constants match the value of the corresponding fields in the message that is at the head of the channel. Again, nothing bad will happen if a statement happens to be non-executable. The process trying to execute it will be delayed until the statement, or, more likely, an alternative statement, becomes executable.

Here is an example that uses some of the mechanisms introduced so far.

```
proctype A(chan q1)
{       chan q2;
        q1?q2;
        q2!123
}

proctype B(chan qforb)
{       int x;
        qforb?x;
        printf("x = %d\n", x)
}
```

```
init {
        chan qname = [1] of { chan };
        chan qforb = [1] of { int };
        run A(qname);
        run B(qforb);
        qname!qforb
}
```
The value printed will be 123.

A predefined function `len(qname)` returns the number of messages currently stored in channel `qname`. Note that if `len` is used as a statement, rather than on the right hand side of an assignment, it will be unexecutable if the channel is empty: it returns a zero result, which by definition means that the statement is temporarily unexecutable. Composite conditions such as

```
        (qname?var == 0)       /* syntax error */
```
or
```
        (a > b && qname!123)   /* syntax error */
```
are invalid in `Promela` (note that these conditions can not be evaluated without side-effects). For a receive statement there is an alternative, using square brackets around the clause behind the question mark.
```
        qname?[ack,var]
```
is evaluated as a condition. It returns 1 if the corresponding receive statement
```
        qname?ack,var
```
is executable, i.e., if there is indeed a message `ack` at the head of the channel. It returns 0 otherwise. In neither case has the evaluation of a statement such as
```
        qname?[ack,var]
```
any side-effects: the receive is evaluated, not executed.

Note carefully that in non-atomic sequences of two statements such as

```
        (len(qname) < MAX) -> qname!msgtype
```
or
```
        qname?[msgtype] -> qname?msgtype
```
the second statement is not *necessarily* executable after the first one has been executed. There may be race conditions if access to the channels is shared between several processes. In the first case another process can send a message to channel `qname` just after this process determined that the channel was not full. In the second case, the other process can steal away the message just after our process determined its presence.

Rendez-Vous Communication

So far we have talked about asynchronous communication between processes via message channels, declared in statements such as

```
        chan qname = [N] of { byte }
```
where `N` is a positive constant that defines the buffer size. A logical extension is to allow for the declaration
```
        chan port = [0] of { byte }
```
to define a rendezvous port that can pass single byte messages. The channel size is zero, that is, the channel `port` can pass, but can not store messages. Message interactions via such rendezvous ports are by definition synchronous. Consider the following example.

```
#define msgtype 33

chan name = [0] of { byte, byte };

proctype A()
{       name!msgtype(124);
        name!msgtype(121)
}

proctype B()
{       byte state;
        name?msgtype(state)
}

init
{       atomic { run A(); run B() }
}
```

Channel `name` is a global rendezvous port. The two processes will synchronously execute their first statement: a handshake on message `msgtype` and a transfer of the value 124 to local variable `state`. The second statement in process `A` will be unexecutable, because there is no matching receive operation in process `B`.

If the channel `name` is defined with a non-zero buffer capacity, the behavior is different. If the buffer size is at least 2, the process of type `A` can complete its execution, before its peer even starts. If the buffer size is 1, the sequence of events is as follows. The process of type `A`can complete its first send action, but it blocks on the second, because the channel is now filled to capacity. The process of type `B` can then retrieve the first message and complete. At this point `A` becomes executable again and completes, leaving its last message as a residual in the channel.

Rendez-vous communication is binary: only two processes, a sender and a receiver, can be synchronized in a rendezvous handshake. We will see an example of a way to exploit this to build a semaphore below. But first, let us introduce a few more control flow structures that may be useful.

Control Flow

Between the lines, we have already introduced three ways of defining control flow: concatenation of statements within a process, parallel execution of processes, and atomic sequences. There are three other control flow constructs in `Promela` to be discussed. They are case selection, repetition, and unconditional jumps.

Case Selection

The simplest construct is the selection structure. Using the relative values of two variables `a` and `b` to choose between two options, for instance, we can write:

```
if
:: (a != b) -> option1
:: (a == b) -> option2
fi
```

The selection structure contains two execution sequences, each preceded by a double colon. Only one sequence from the list will be executed. A sequence can be selected only if its first statement is executable. The first statement is therefore called a *guard*.

In the above example the guards are mutually exclusive, but they need not be. If more than one guard is executable, one of the corresponding sequences is selected nondeterministically. If all guards are unexecutable the process will block until at least one of them can be selected. There is no restriction on the type of statements that can be used as a guard. The following example, for instance, uses input statements.

```
#define a 1
#define b 2

chan ch = [1] of { byte };

proctype A()
{       ch!a
}

proctype B()
{       ch!b
}

proctype C()
{       if
        :: ch?a
        :: ch?b
        fi
}

init
{       atomic { run A(); run B(); run C() }
}
```

The example defines three processes and one channel. The first option in the selection structure of the process of type c is executable if the channel contains a message a, where a is a constant with value 1, defined in a macro definition at the start of the program. The second option is executable if it contains a message b, where, similarly, b is a constant. Which message will be available depends on the unknown relative speeds of the processes.

A process of the following type will either increment or decrement the value of variable count once.

```
byte count;

proctype counter()
{
        if
        :: count = count + 1
        :: count = count - 1
        fi
}
```

Repetition

A logical extension of the selection structure is the [repetition structure](#). We can modify the above program as follows, to obtain a cyclic program that randomly changes the value of the variable up or down.

```
byte count;

proctype counter()
{
        do
        :: count = count + 1
        :: count = count - 1
        :: (count == 0) -> break
        od
}
```

Only one option can be selected for execution at a time. After the option completes, the execution of the structure is repeated. The normal way to terminate the repetition structure is with a `break` statement. In the example, the loop can be broken when the count reaches zero. Note, however, that it need not terminate since the other two options always remain executable. To force termination when the counter reaches zero, we could modify the program as follows.

```
proctype counter()
{
        do
        :: (count != 0) ->
                if
                :: count = count + 1
                :: count = count - 1
                fi
        :: (count == 0) -> break
        od
}
```

Unconditional Jumps

Another way to break the loop is with an unconditional jump: the infamous `goto` statement. This is illustrated in the following implementation of Euclid's algorithm for finding the greatest common divisor of two non-zero, positive numbers:

```
proctype Euclid(int x, y)
{
        do
        :: (x >  y) -> x = x - y
        :: (x <  y) -> y = y - x
        :: (x == y) -> goto done
        od;
done:
        skip
}
```

The `goto` in this example jumps to a label named done. A label can only appear before a statement. Above we want to jump to the end of the program. In this case a dummy statement `skip` is useful: it is a place holder that is always executable and has no effect. The `goto` is also always executable.

The following example specifies a filter that receives messages from a channel `in` and divides them over two channels large and small depending on the values attached. The constant `N` is defined to be 128 and size is defined to be 16 in the two macro definitions.

```
#define N      128
#define size   16

chan in    = [size] of { short };
chan large = [size] of { short };
chan small = [size] of { short };

proctype split()
{      short cargo;

       do
       :: in?cargo ->
               if
               :: (cargo >= N) ->
                       large!cargo
               :: (cargo <  N) ->
                       small!cargo
               fi
       od
}

init
{      run split()
}
```

A process type that merges the two streams back into one, most likely in a different order, and writes it back into the channel `in` could be specified as follows.

```
proctype merge()
{      short cargo;

       do
       ::      if
               :: large?cargo
               :: small?cargo
               fi;
               in!cargo
       od
}
```

If we now modify the `init` process as follows, the split and merge processes could busily perform their duties forever on.

```
init
{      in!345; in!12; in!6777;
       in!32;  in!0;
       run split();
       run merge()
}
```

As a final example, consider the following implementation of a Dijkstra semaphore, using binary rendezvous communication.

```
#define p      0
#define v      1

chan sema = [0] of { bit };
```

```
proctype dijkstra()
{       byte count = 1;

        do
        :: (count == 1) ->
                sema!p; count = 0
        :: (count == 0) ->
                sema?v; count = 1
        od
}

proctype user()
{       do
        :: sema?p;
           /*     critical section */
           sema!v;
           /* non-critical section */
        od
}

init
{       run dijkstra();
        run user();
        run user();
        run user()
}
```
The semaphore guarantees that only one of the user processes can enter its critical section at a time. It does not necessarily prevent the monopolization of the access to the critical section by one of the processes.

Modeling Procedures and Recursion

Procedures can be modeled as processes, even recursive ones. The return value can be passed back to the calling process via a global variable, or via a message. The following program illustrates this.

```
proctype fact(int n; chan p)
{       chan child = [1] of { int };
        int result;

        if
        :: (n <= 1) -> p!1
        :: (n >= 2) ->
                run fact(n-1, child);
                child?result;
                p!n*result
        fi
}
init
{       chan child = [1] of { int };
        int result;

        run fact(7, child);
        child?result;
        printf("result: %d\n", result)
}
```
The process *fact(n, p)* recursively calculates the factorial of *n* , communicating the result via a message to its parent process *p*.

Timeouts

We have already discussed two types of statement with a predefined meaning in `Promela`: `skip`, and `break`. Another predefined statement is `timeout`. The `timeout` models a special condition that allows a process to abort the waiting for a condition that may never become true, e.g. an input from an empty channel. The timeout keyword is a modeling feature in `Promela` that provides an escape from a hang state. The timeout condition becomes true only when no other statements within the distributed system is executable. Note that we deliberately abstract from absolute timing considerations, which is crucial in verification work, and we do not specify how the timeout should be implemented. A simple example is the following process that will send a reset message to a channel named *guard* whenever the system comes to a standstill.

```
proctype watchdog()
{
        do
        :: timeout -> guard!reset
        od
}
```

Assertions

Another important language construct in `Promela` that needs little explanation is the `assert` statement. Statements of the form

```
assert(any_boolean_condition)
```
are always executable. If the boolean condition specified holds, the statement has no effect. If, however, the condition does not necessarily hold, the statement will produce an error report during verifications with *Spin*.

Advanced Usage

The modeling language has a few features that specifically address the verification aspects. It shows up in the way labels are used, in the semantics of the `Promela timeout` statement, and in the usage of statements such as `assert` that we discuss next.

End-State Labels

When `Promela` is used as a verification language the user must be able to make very specific assertions about the behavior that is being modeled. In particular, if a `Promela` is checked for the presence of deadlocks, the verifier must be able to distinguish a normal *end state* from an abnormal one.

A normal end state could be a state in which every `Promela` process that was instantiated has properly reached the end of the defining program body, and all message channels are empty. But, not all `Promela` process are, of course, meant to reach the end of their

program body. Some may very well linger in an `IDLE` state, or they may sit patiently in a loop ready to spring into action when new input arrives.

To make it clear to the verifier that these alternate end states are legal, and do not constitute a deadlock, a `Promela` model can use end state labels. For instance, if by adding a label to the process type `dijkstra()` we discussed earlier:

```
proctype dijkstra()
{       byte count = 1;

end:    do
        :: (count == 1) ->
               sema!p; count = 0
        :: (count == 0) ->
               sema?v; count = 1
        od
}
```

we indicate that it is not an error if, at the end of an execution sequence, a process of type `dijkstra()` has not reached its closing curly brace, but waits in the loop. Of course, such a state could still be part of a deadlock state, but if so, it is not caused by this particular process. (It will still be reported if any one of the other processes in not in a valid end-state).

There may be more than one end state label per verification model. If so, all labels that occur within the same process body must be unique. The rule is that every label name that *starts* with the three character sequence `"end"` is an endstate label. So it is perfectly valid to use variations such as `enddne`, `end0`, `end_appel`, etc.

Progress-State Labels

In the same spirit as the end state labels, the user can also define *progress state* labels. In this case, a progress state labels will mark a state that *must* be executed for the protocol to make progress. Any infinite cycle in the protocol execution that does not pass through at least one of these progress states, is a potential starvation loop. In the `dijkstra` example, for instance, we can label the successful passing of a semaphore test as ``progress'' and ask a verifier to make sure that there is no cycle in the protocol execution where at least one process succeeds in passing the semaphore guard.

```
proctype dijkstra()
{       byte count = 1;

end:    do
        :: (count == 1) ->
progress:      sema!p; count = 0
        :: (count == 0) ->
               sema?v; count = 1
        od
}
```

Note that a label cannot be added immediately before or immediately after the "::" symbol, so you may have to be more creative to find a good place to put it.
If more than one state carries a progress label, variations with a common prefix are again valid: `progress0`, `progress_foo`, etc.

All analyzers generated by *Spin* with the `-a` flag have a runtime option (after compilation) named `-l`. Invoking the generated analyzer with that flag will cause a fast search for non-progress loops, instead of the default search for deadlocks. The search takes about twice as long (and uses twice as much memory) as the default search for deadlocks. (A considerable improvement over standard methods that are based on the analysis of strongly connected components.)

Message Type Definitions

We have seen how variables are declared and how constants can be defined using C-style macros. `Promela` also allows for message type definitions that look as follows:

```
mtype = { ack, nak, err, next, accept }
```
This is a preferred way of specifying the message types since it abstracts from the specific values to be used, and it makes the names of the constants available to an implementation, which can improve error reporting.

By using the `mtype` keyword in channel declarations, the corresponding message field will always be interpreted symbolically, instead of numerically. For instance:

```
chan q = [4] of { mtype, mtype, bit, short };
```

Pseudo Statements

We have now discussed all the basic types of statements defined in `Promela`: assignments, conditions, send and receive, `assert`, `timeout`, `goto`, `break` and `skip`. Note that `chan`, `len` and `run` are not really statements but unary operators that can be used in conditions and assignments.

The `skip` statement was mentioned in passing as a statement that can be a useful filler to satisfy syntax requirements, but that really has no effect. It is formally not part of the language but a *pseudo-statement*, merely a synonym of another statement with the same effect: a simple condition of a constant value (1). In the same spirit other pseudo-statements could be defined (but are not), such as `block` or `hang`, as equivalents of (0), and `halt`, as an equivalent of `assert(0)`.. Another pseudo-statement is `else`, that can be used as the initial statement of the last option sequence in a selection or iteration.

```
if
:: a > b -> ...
:: else -> ...
fi
```
The `else` is only executable (true) if all other options in the same selection are not executable.

Example

Here is a simple example of a (flawed) protocol, modeled in `Promela`.

```
mtype = { ack, nak, err, next, accept };

proctype transfer(chan in,out,chin,chout)
{       byte o, i;

        in?next(o);

        do
        :: chin?nak(i) ->
                        out!accept(i);
                        chout!ack(o)

        :: chin?ack(i) ->
                        out!accept(i);
                        in?next(o);
                        chout!ack(o)

        :: chin?err(i) ->
                        chout!nak(o)
        od
}

init
{       chan AtoB = [1] of { mtype, byte };
        chan BtoA = [1] of { mtype, byte };

        chan Ain  = [2] of { mtype, byte };
        chan Bin  = [2] of { mtype, byte };

        chan Aout = [2] of { mtype, byte };
        chan Bout = [2] of { mtype, byte };

        atomic {
          run transfer(Ain,Aout, AtoB,BtoA);
          run transfer(Bin,Bout, BtoA,AtoB)
        };

        AtoB!err(0)
}
```

The channels `Ain` and `Bin` are to be filled with token messages of type `next` and arbitrary values (e.g. ASCII character values) by unspecified background processes: the users of the transfer service. Similarly, these user processes can read received data from the channels`Aout` and `Bout`. The channels and processes are initialized in a single atomic statement, and started with the dummy `err` message.

Spin (see also Spin.html)

Given a model system specified in `Promela`, *Spin* can either perform random simulations of the system's execution or it can generate a C program that performs a fast exhaustive verification of the system state space. The verifier can check, for instance, if user specified system invariants may be violated during a protocol's execution.

If *Spin* is invoked without any options it performs a random simulation. With option `-n` `N` the seed for the simulation is set explicitly to the integer value N.

A group of options `pglrs` can be used to set the desired level of information that the user wants about the simulation run. Every line of output normally contains a reference to the source line in the specification that caused it.

`-p` Shows the state changes of the `Promela` processes at every time step.

`-g` Shows the current value of global variables at every time step.

`-l` Shows the current value of local variables, after the process that owns them has changed state. It is best used in combination with option `-p`.

`-r` Shows all message receive events. It shows the process performing the receive, its name and number, the source line number, the message parameter number (there is one line for each parameter), the message type and the message channel number and name.

`-s` Shows all message send events.

Spin understands three other options (`-a`, `-m`, and `-t`):

`-a` Generates a protocol specific analyzer. The output is written into a set of C files, named `pan.[cbhmt]`, that can be compiled to produce the analyzer (which is then executed to perform the analysis). To guarantee an exhaustive exploration of the state space, the program can be compiled simply as

```
$ gcc -o pan pan.c
```
For larger systems this may, however, exhaust the available memory on the machine used. Large to very large systems can still be analyzed by using a memory efficient bit state space method by
```
$ gcc -DBITSTATE -o pan pan.c
```
An indication of the coverage of such a search can be derived from the *hash factor* (see below). The generated executable analyzer, named `run` above, has its own set of options that can be seen by typing "`./pan -?`" (see also below in *Using the Analyzer*).

`-m` can be used to change the default semantics of send actions. Normally, a send operation is only executable if the target channel is non-full. This imposes an implicit synchronization that can not always be justified. Option `-m` causes send actions to be always executable. Messages sent to a channel that is full are then dropped. If this option is combined with `-a` the semantics of send in the analyzers generated is similarly altered, and the verifications will take the effects of this type of message loss into consideration.

`-t` is a trail-hunting option. If the analyzer finds a violation of an assertion, a deadlock or an unspecified reception, it writes an error trail into a file named `pan.trail`. The trail can be inspected in detail by invoking *Spin* with the `-t` option. In combination with the options`pglrs` different views of the error sequence are then easily obtained.

For brevity, other options of *Spin* are not discussed here. For details see [5]. For a hint of their purpose, see ``Digging Deeper'' at the end of this manual.

The Simulator

Consider the following example protocol, that we will store in a file named `lynch`.

```
1   #define MIN      9
2   #define MAX      12
3   #define FILL     99
4
5   mtype = { ack, nak, err }
6
7   proctype transfer(chan chin, chout)
8   { byte o, i, last_i=MIN;
9
10    o = MIN+1;
11    do
12    :: chin?nak(i) ->
13          assert(i == last_i+1);
14          chout!ack(o)
15    :: chin?ack(i) ->
16          if
17          :: (o <  MAX) -> o = o+1
18          :: (o >= MAX) -> o = FILL
19          fi;
20          chout!ack(o)
21    :: chin?err(i) ->
22          chout!nak(o)
23    od
24  }
25
26  proctype channel(chan in, out)
27  { byte md, mt;
28    do
29    :: in?mt,md ->
30          if
31          :: out!mt,md
32          :: out!err,0
33          fi
34    od
35  }
36
37  init
38  { chan AtoB = [1] of { mtype, byte };
39    chan BtoC = [1] of { mtype, byte };
40    chan CtoA = [1] of { mtype, byte };
41    atomic {
42          run transfer(AtoB, BtoC);
43          run channel(BtoC, CtoA);
44          run transfer(CtoA, AtoB)
45    };
46    AtoB!err,0;    /* start */
47    0              /* hang */
48  }
```

The protocol uses three message types: `ack`, `nak`, and a special type `err` that is used to model message distortions on the communication channel between the two transfer processes. The behavior of the channel is modeled explicitly with a channel process. There is also an assert statement that claims a, faulty, invariant relation between two local variables in the transfer processes.

Running *Spin* without options gives us a random simulation that will only provide output when execution terminates, or if a `printf` statement is encountered. In this case:

```
$ spin lynch    # no options, not recommended
spin: "lynch" line 13: assertion violated
#processes: 4
proc  3 (transfer)    line 11 (state 15)
proc  2 (channel)     line 28 (state 6)
proc  1 (transfer)    line 13 (state 3)
proc  0 (:init:)      line 48 (state 6)
4 processes created
$
```

There are no `printf`'s in the specification, but we do get some output because the execution stops on an assertion violation. Curious to find out more, we can repeat the run with more verbose output, e.g. printing all receive events. The result of that run is shown in Figure 1. Most output will be self-explanatory.

The above simulation run ends in the same assertion violation. Since the simulation resolves nondeterministic choices in a random manner, this need not always be the case. To force a reproducible run, the option `-n N` can be used. For instance:

```
$ spin -r -n100 lynch
```

will seed the random number generator with the integer value 100 and is guaranteed to produce the same output each time it is executed.

The other options can add still more output to the simulation run, but the amount of text can quickly become overwhelming. An easy solution is to filter the output through *grep*. For instance, if we are only interested in the behavior of the channel process in the above example, we say:

```
$ spin -n100 -r lynch | grep "proc  2"
```

The results are shown in Figure 1.

```
$ spin -r lynch
proc  1 (transfer) line  21, Recv err,0  <- queue 1 (chin)
proc  2 (channel)  line  29, Recv nak,10 <- queue 2 (in)
proc  3 (transfer) line  12, Recv nak,10 <- queue 3 (chin)
proc  1 (transfer) line  15, Recv ack,10 <- queue 1 (chin)
...
proc  1 (transfer) line  15, Recv ack,12 <- queue 1 (chin)
proc  2 (channel)  line  29, Recv ack,99 <- queue 2 (in)
proc  3 (transfer) line  15, Recv ack,99 <- queue 3 (chin)
proc  1 (transfer) line  15, Recv ack,99 <- queue 1 (chin)
proc  2 (channel)  line  29, Recv ack,99 <- queue 2 (in)
proc  3 (transfer) line  21, Recv err,0  <- queue 3 (chin)
proc  1 (transfer) line  12, Recv nak,99 <- queue 1 (chin)
spin: "lynch" line 13: assertion violated
#processes: 4
proc  3 (transfer) line 11 (state 15)
proc  2 (channel)  line 28 (state 6)
proc  1 (transfer) line 13 (state 3)
proc  0 (:init:)   line 48 (state 6)
4 processes created
$ spin -n100 -r lynch | grep "proc  2"
proc  2 (channel) line 29, Recv nak,10 <- queue 2 (in)
proc  2 (channel) line 29, Recv ack,11 <- queue 2 (in)
proc  2 (channel) line 29, Recv ack,12 <- queue 2 (in)
proc  2 (channel) line 28 (state 6)

Figure 1.  Simulation Run Output
```

A good way to start simulation runs like this is to use option `-c` (new in version 3.0), which produces the following output:

```
$ spin -c lynch
proc 0 = :init:
proc 1 = transfer
proc 2 = channel
proc 3 = transfer
q\p  0   1    2     3
  1  AtoB!err,0
  1  .   chin?err,0
  2  .   chout!nak,10
  2  .   .    in?nak,10
  3  .   .    out!err,0
  3  .   .    .     chin?err,0
  1  .   .    .     chout!nak,10
  1  .   chin?nak,10
  2  .   chout!ack,10
  2  .   .    in?ack,10
  3  .   .    out!ack,10
  3  .   .    .     chin?ack,10
  1  .   .    .     chout!ack,11
  1  .   chin?ack,11
  2  .   chout!ack,11
  2  .   .    in?ack,11
  3  .   .    out!ack,11
  3  .   .    .     chin?ack,11
  1  .   .    .     chout!ack,12
  1  .   chin?ack,12
  2  .   chout!ack,12
  2  .   .    in?ack,12
  3  .   .    out!ack,12
  3  .   .    .     chin?ack,12
  1  .   .    .     chout!ack,99
  1  .   chin?ack,99
  2  .   chout!ack,99
  2  .   .    in?ack,99
  3  .   .    out!err,0
  3  .   .    .     chin?err,0
  1  .   .    .     chout!nak,99
  1  .   chin?nak,99
spin: line  13 "lynch", Error: assertion violated
-------------
final state:
-------------
#processes: 4
 79:    proc  3 (transfer) line  11 "lynch" (state 15)
 79:    proc  2 (channel) line  28 "lynch" (state 6)
 79:    proc  1 (transfer) line  13 "lynch" (state 3)
 79:    proc  0 (:init:) line  47 "lynch" (state 6)
4 processes created
```

The first column gives the channel id numbers, the first row gives the process id numbers involved in message transfers plotted in the remained of the table.

The Analyzer (see also Pan.html and Roadmap.html)

The simulation runs can be useful in quick debugging of new designs, but by simulation alone we can not prove that the system is really error free. A verification of even very large models can be performed with the `-a` and `-t` options of *Spin*.

An exhaustive state space searching program for a protocol model is generated as follows, producing five files, named *pan.[bchmt]*.

```
$ spin -a lynch
$ wc pan.[bchmt]
      99     285    1893 pan.b
    3158   10208   70337 pan.c
     356    1238    7786 pan.h
     216     903    6045 pan.m
     575    2099   14017 pan.t
    4404   14733  100078 total
```
The details are none too interesting: *pan.c* contains most of the C code for the analysis of the protocol. File *pan.t* contains a transition matrix that encodes the protocol control flow; *pan.b* and *pan.m* contain C code for forward and backward transitions and *pan.h* is a general header file. The program can be compiled in several ways, e.g., with a full state space or with a bit state space.

Exhaustive Search

The best method, that works up to system state spaces of roughly 100,000 states, is to use the default compilation of the program:

```
$ gcc -o pan pan.c
```
The executable program *pan* can now be executed to perform the verification. The verification is truly exhaustive: it tests all possible event sequences in all possible orders. It should, of course, find the same assertion violation.
```
$ ./pan
assertion violated (i == last_i + 1))
pan: aborted
pan: wrote pan.trail
search interrupted
vector 64 byte, depth reached 56
      61 states, stored
       5 states, linked
       1 states, matched
hash conflicts: 0 (resolved)
(size 2^18 states, stack frames: 0/5)
```
(The output format of the more recent versions of *spin* is more elaborate, but it includes the same information.)

The first line of the output announces the assertion violation and attempts to give a first indication of the invariant that was violated. The violation was found after 61 states had been generated. Hash "conflicts" gives the number of hash collisions that happened during access to the state space. As indicated, all collissions are resolved in full search mode, since all states are placed in a linked list. The most relevant piece of output in this case, however, is on the third line which tells us that a trail file was created that can be used in combination with the simulator to recreate the error sequence. We can now say, for instance

```
$ spin -t -r lynch | grep "proc  2"
```
to determine the cause of the error. Note carefully that the verifier is guaranteed to find the assertion violation if it is feasible. If an exhaustive search does not report such a violation, it is certain that *no* execution execution sequence exists that can violate the assertion.

The executable analyzer that is generated comes with a modest number of options that can be checked as follows

```
$ ./pan --
-cN stop at Nth error (default=1)
-l  find non-progress cycles  # when compiled with -DNP
-a  find acceptance cycles    # when not compiled with -DNP
-mN max depth N (default=10k)
-wN hash table of 2^N entries (default=18)
 ...etc.
```

Using a zero as an argument to the first option forces the state space search to continue, even if errors are found. An overview of unexecutable (unreachable) code is given with every complete run: either the default run if it did not find any errors, or the run with option -c0. In this case the output is:

```
$ ./pan -c0
assertion violated (i == (last_i + 1))

vector 64 byte, depth reached 60, errors: 5
     165 states, stored
       5 states, linked
      26 states, matched
hash conflicts: 1 (resolved)
(size 2^18 states, stack frames: 0/6)

unreached code :init: (proc 0):
        reached all 9 states
unreached code channel (proc 1):
         line 35 (state 9),
         reached: 8 of 9 states
unreached code transfer (proc 2):
         line 24 (state 18),
         reached: 17 of 18 states
```

There were five assertion violations, and some 165 unique system states were generated. Each state description (the *vector size*) took up 64 bytes of memory; the longest non-cyclic execution sequence was 60. There is one unreachable state both in the channel process and in the transfer process. In both cases the unreachable state is the control flow point just after the do-loop in each process. Note that both loops are indeed meant to be non-terminating.

The -l option will cause the analyzer to search for non-progress loops rather than deadlocks or assertion violations. The option is explained in the section on ``More Advanced Usage."

The executable analyzer has two other options. By default the search depth is restricted to a rather arbitrary 10,000 steps. If the depth limit is reached, the search is truncated, making the verification less than exhaustive. To make certain that the search is exhaustive, make sure that the "depth reached" notice is within the maximum search depth, and if not, repeat the analysis with an explicit -m argument.

The -m option can of course also be used to truncate the search explicitly, in an effort to find the shortest possible execution sequence that violates a given assertion. Such a truncated search, however, is not guaranteed to find every possible violation, even within the search depth.

The last option `-w N` can only affect the run time, not the scope, of an analysis with a full state space. This "hash table width" should normally be set equal to, or preferably higher than, the logarithm of the expected number of unique system states generated by the analyzer. (If it is set too low, the number of hash collisions will increase and slow down the search.) The default N of 18 handles up to 262,144 system states, which should suffice for almost all applications of a full state space analysis.

Bit State Space Analyses

It can easily be calculated what the memory requirements of an analysis with a full state space are [4]. If, as in the example we have used, the protocol requires 64 bytes of memory to encode one system state, and we have a total of 2MB of memory available for the search, we can store up to 32,768 states. The analysis fails if there are more reachable states in the system state space. So far, *spin* is the *only* verification system that can avoid this trap. All other existing automated verification system (irrespective on which formalism they are based) simply run out of memory and abort their analysis without returning a useful answer to the user.

The coverage of a conventional analysis goes down rapidly when the memory limit is hit, i.e. if there are twice as many states in the full state space than we can store, the effective coverage of the search is only 50% and so on. *spin* does substantially better in those cases by using the bit state space storage method [4]. The bit state space can be included by compiling the analyzer as follows:

```
$ gcc -DBITSTATE -o pan pan.c
```
The analyzer compiled in this way should of course find the same assertion violation again:
```
$ ./pan
assertion violated (i == ((last_i + 1))
pan: aborted
pan: wrote pan.trail
search interrupted
vector 64 byte, depth reached 56
     61 states, stored
      5 states, linked
      1 states, matched
hash factor: 67650.064516
(size 2^22 states, stack frames: 0/5)
$
```
In fact, for small to medium size problems there is very little difference between the full state space method and the bit state space method (with the exception that the latter is somewhat faster and uses substantially less memory). The big difference comes for larger problems. The last two lines in the output are useful in estimating the *coverage* of a large run. The maximum number of states that the bit state space can accommodate is written on the last line (here 2^{22} bytes or about 32 million bits = states). The line above it gives the *hash factor*: roughly equal to the maximum number of states divided by the actual number of states. A large hash factor (larger than 100) means, with high reliability, a coverage of 99% or 100%. As the hash factor approaches 1 the coverage approaches 0%.

Note carefully that the analyzer realizes a partial coverage *only* in cases where traditional verifiers are either unable to perform a search, or realize a far smaller coverage.

In *no* case will `Spin` produce an answer that is less reliable than that produced by other automated verification systems (quite on the contrary).

The object of a bit state verification is to achieve a hash factor larger than 100 by allocating the maximum amount of memory for the bit state space. For the best result obtainable: use the `-w N` option to size the state space to precisely the amount of real (not virtual) memory available on your machine. By default, *N* is 22, corresponding to a state space of 4MB. For example, if your machine has 128MB of real memory, you can use `-w27` to analyze systems with up to a billion reachable states.

Summary

In the first part of this manual we have introduced a notation for modeling concurrent systems, including but not limited to asynchronous data communication protocols, in a language named `Promela`. The language has several unusual features. All communication between processes takes place via either messages or shared variables. Both synchronous and asynchronous communication are modeled as two special cases of a general message passing mechanism. Every statement in `Promela` can potentially model delay: it is either executable or not, in most cases depending on the state of the environment of the running process. Process interaction and process coordination is thus at the very basis of the language. More about the design of `Promela`, of the verifier `Spin`, and its range of applications, can be found in [5].

`Promela` is deliberately a verification modeling language, not a programming language. There are, for instance, no elaborate abstract data types, or more than a few basic types of variable. A verification model is an abstraction of a protocol implementation. The abstraction maintains the essentials of the process interactions, so that it can be studied in isolation. It suppresses implementation and programming detail.

The syntax of `Promela` expressions, declarations, and assignments is loosely based on the language `C`[6]. The language was influenced significantly by the ``guarded command languages'' of E.W. Dijkstra [1] and C.A.R. Hoare [3]. There are, however, important differences. Dijkstra's language had no primitives for process interaction. Hoare's language was based exclusively on synchronous communication. Also in Hoare's language, the type of statements that could appear in the guards of an option was restricted. The semantics of the selection and cycling statements in `Promela` is also rather different from other guarded command languages: the statements are not aborted when all guards are false but they block: thus providing the required synchronization.

With minimal effort `Spin` allows the user to generate sophisticated analyzers from `Promela` verification models. Both the `Spin` software itself, and the analyzers it can generate, are written in ANSI C and are portable across *Unix* systems. They can be scaled to fully exploit the physical limitations of the host computer, and deliver within those limits the best possible analyses that can be realized with the current state of the art in protocol analysis.

Appendix: Building a Verification Suite

The first order of business in using *Spin* for a verification is the construction of a faithful model in `Promela` of the problem at hand. The language is deliberately kept small. The purpose of the modeling is to extract those aspects of the system that are relevant to the coordination problem being studied. All other details are suppressed. Formally: the model is a reduction of the system that needs to be equivalent to the full system only with respect to the properties that are being verified. Once a model has been constructed, it becomes the basis for the construction of a series of, what we may call, ``verification suites'' that are used to verify its properties. To build a verification suite we can prime the model with assertions. The assertions can formalize invariant relations about the values of variables or about allowable sequences of events in the model.

As a first example we take the following solution to the mutual exclusion problem, discussed earlier, published in 1966 by H. Hyman in the Communications of the ACM. It was listed, in pseudo Algol, as follows.

```
 1 Boolean array b(0;1) integer k, i,
 2 comment process i, with i either 0 or 1, and k = 1-i;
 3 C0: b(i) := false;
 4 C1: if k != i then begin;
 5 C2: if not b(1-i) then go to C2;
 6     else k := i; go to C1 end;
 7     else critical section;
 8     b(i) := true;
 9     remainder of program;
10     go to C0;
11     end
```

The solution, as Dekker's earlier solution, is for two processes, numbered 0 and 1. Suppose we wanted to prove that Hyman's solution truly guaranteed mutually exclusive access to the critical section. Our first task is to build a model of the solution in `Promela`. While we're at it, we can pick some more useful names for the variables that are used.

```
 1  bool want[2];    /* Bool array b */
 2  bool turn;       /* integer    k */
 3
 4  proctype P(bool i)
 5  {
 6    want[i] = 1;
 7    do
 8    :: (turn != i) ->
 9            (!want[1-i]);
10            turn = i
11    :: (turn == i) ->
12            break
13    od;
14    skip; /* critical section */
15    want[i] = 0
16  }
17
18  init { run P(0); run P(1) }
```

We can generate, compile, and run a verifier for this model, to see if there are any major problems, such as a global system deadlock.

```
        $ spin -a hyman0
        $ gcc -o pan pan.c
```

```
$ ./pan
full statespace search for:
assertion violations and invalid endstates
vector 20 byte, depth reached 19, errors: 0
      79 states, stored
       0 states, linked
      38 states, matched      total: 117
hash conflicts: 4 (resolved)
(size 2^18 states, stack frames: 3/0)


unreached code _init (proc 0):
      reached all 3 states
unreached code P (proc 1):
      reached all 12 states
```

The model passes this first test. What we are really interested in, however, is if the algorithm guarantees mutual exclusion. There are several ways to proceed. The simplest is to just add enough information to the model that we can express the correctness requirement in a Promela assertion.

```
1   bool want[2];
2   bool turn;
3   byte cnt;
4
5   proctype P(bool i)
6   {
7     want[i] = 1;
8     do
9     :: (turn != i) ->
10            (!want[1-i]);
11            turn = i
12    :: (turn == i) ->
13            break
14    od;
15    skip; /* critical section */
16    cnt = cnt+1;
17    assert(cnt == 1);
18    cnt = cnt-1;
19    want[i] = 0
20  }
21
22  init { run P(0); run P(1) }
```

We have added a global variable cnt that is incremented upon each access to the critical section, and decremented upon each exit from it. The maximum value that this variable should ever have is 1, and it can only have this value when a process is inside the critical section.

```
$ spin -a hyman1
$ gcc -o pan pan.c
$ ./pan
assertion violated (cnt==1)
pan: aborted (at depth 15)
pan: wrote pan.trail
full statespace search for:
assertion violations and invalid endstates
search was not completed
vector 20 byte, depth reached 25, errors: 1
     123 states, stored
       0 states, linked
      55 states, matched      total: 178
hash conflicts: 42 (resolved)
(size 2^18 states, stack frames: 3/0)
```

The verifier claims that the assertion can be violated. We can use the error trail to check it with *Spin*'s -t option:

```
$ spin -t -p hyman1
proc  0 (_init) line 24 (state 2)
proc  0 (_init) line 24 (state 3)
proc  2 (P)     line 8 (state 7)
proc  2 (P)     line 9 (state 2)
proc  2 (P)     line 10 (state 3)
proc  2 (P)     line 11 (state 4)
proc  1 (P)     line 8 (state 7)
proc  1 (P)     line 12 (state 5)
proc  1 (P)     line 15 (state 10)
proc  2 (P)     line 8 (state 7)
proc  2 (P)     line 12 (state 5)
proc  2 (P)     line 15 (state 10)
proc  2 (P)     line 16 (state 11)
proc  2 (P)     line 17 (state 12)
proc  2 (P)     line 18 (state 13)
proc  1 (P)     line 16 (state 11)
proc  1 (P)     line 17 (state 12)
spin: "hyman1" line 17: assertion violated
step 17, #processes: 3
                    want[0] = 1
                    _p[0] = 12
                    turn[0] = 1
                    cnt[0] = 2

proc  2 (P)     line 18 (state 13)
proc  1 (P)     line 17 (state 12)
proc  0 (_init) line 24 (state 3)
3 processes created
```

Here is another way to catch the error. We again lace the model with the information that
will allow us to count the number of processes in the critical section.

```
1   bool want[2];
2   bool turn;
3   byte cnt;
4
5   proctype P(bool i)
6   {
7     want[i] = 1;
8     do
9     :: (turn != i) ->
10           (!want[1-i]);
11           turn = i
12    :: (turn == i) ->
13           break
14    od;
15    cnt = cnt+1;
16    skip;   /* critical section */
17    cnt = cnt-1;
18    want[i] = 0
19  }
20
21  proctype monitor()
22  {
23    assert(cnt == 0 || cnt == 1)
24  }
25
26  init {
27    run P(0); run P(1); run monitor()
28  }
```

The invariant condition on the value of counter cnt is now place in a separate
process monitor() (the name is immaterial). The extra process runs along with the two
others. It will always terminate in one step, but it could execute that step at *any* time.
The systems modeled by Promela and verified by *Spin* are completely asynchronous.

That means that the verification of *Spin* take into account *all* possible relative timings of the three processes. In a full verification, the assertion therefore can be evaluated at any time during the lifetime of the other two processes. If the verifier reports that it is not violated we can indeed conclude that there is no execution sequence at all (no way to select relative speeds for the three processes) in which the assertion can be violated. The setup with the monitor process is therefore an elegant way to check the validity of a system invariant. The verification produces:

```
$ spin -a hyman2
$ gcc -o pan pan.c
$ ./pan
assertion violated ((cnt==0)||(cnt==1))
pan: aborted (at depth 15)
pan: wrote pan.trail
full statespace search for:
assertion violations and invalid endstates
search was not completed
vector 24 byte, depth reached 26, errors: 1
      368 states, stored
        0 states, linked
      379 states, matched      total: 747
hash conflicts: 180 (resolved)
(size 2^18 states, stack frames: 4/0)
```

Because of the extra interleaving of the two processes with a third monitor, the number of system states that had to be searched has increased, but the error is again correctly reported.

Another Example

Not always can a correctness requirement be cast in terms of a global system invariant. Here is an example that illustrates this. It is a simple alternating bit protocol, modeling the possibility of message loss, and distortion, and extended with negative acknowledgements.

```
 1  #define MAX      5
 2
 3  mtype = { mesg, ack, nak, err };
 4
 5  proctype sender(chan in, out)
 6  { byte o, s, r;
 7
 8    o=MAX-1;
 9    do
10    :: o = (o+1)%MAX; /* next msg */
11  again: if
12         :: out!mesg(o,s) /* send */
13         :: out!err(0,0)  /* distort */
14         :: skip          /* or lose */
15         fi;
16         if
17         :: timeout    -> goto again
18         :: in?err(0,0) -> goto again
19         :: in?nak(r,0) -> goto again
20         :: in?ack(r,0) ->
21           if
22           :: (r == s) -> goto progress
23           :: (r != s) -> goto again
24           fi
25         fi;
26  progress: s = 1-s /* toggle seqno */
```

```
27   od
28  }
29
30  proctype receiver(chan in, out)
31  { byte i;          /* actual input    */
32    byte s;          /* actual seqno    */
33    byte es;         /* expected seqno  */
34    byte ei;         /* expected input  */
35
36    do
37    :: in?mesg(i, s) ->
38            if
39            :: (s == es) ->
40                    assert(i == ei);
41  progress:         es = 1 - es;
42                    ei = (ei + 1)%MAX;
43                    if
44  /* send,   */  :: out!ack(s,0)
45  /* distort */  :: out!err(0,0)
46  /* or lose */  :: skip
47                    fi
48            :: (s != es) ->
49                    if
50  /* send,   */  :: out!nak(s,0)
51  /* distort */  :: out!err(0,0)
52  /* or lose */  :: skip
53                    fi
54            fi
55    :: in?err ->
56            out!nak(s,0)
57    od
58  }
59
60  init {
61    chan s_r = [1] of { mtype,byte,byte };
62    chan r_s = [1] of { mtype,byte,byte };
63    atomic {
64            run sender(r_s, s_r);
65            run receiver(s_r, r_s)
66    }
67  }
```

To test the proposition that this protocol will correctly transfer data, the model has already been primed for the first verification runs. First, the sender is setup to transfer an infinite series of integers as messages, where the value of the integers are incremented modulo MAX. The value of MAX is not really too interesting, as long as it is larger than the range of the sequence numbers in the protocol: in this case 2. We want to verify that data that is sent can only be delivered to the receiver without any deletions or reorderings, despite the possibility of arbitrary message loss. The assertion on line 40 verifies precisely that. Note that if it were ever possible for the protocol to fail to meet the above requirement, the assertion can be violated.

A first verification run reassures us that this is not possible.

```
$ spin -a ABP0
$ gcc -o pan pan.c
$ ./pan
full statespace search for:
assertion violations and invalid endstates
vector 40 byte, depth reached 131, errors: 0
     346 states, stored
       1 states, linked
     125 states, matched      total: 472
```

```
hash conflicts: 17 (resolved)
(size 2^18 states, stack frames: 0/25)

unreached code _init (proc 0):
        reached all 4 states
unreached code receiver (proc 1):
        line 58 (state 24)
        reached: 23 of 24 states
unreached code sender (proc 2):
        line 28 (state 27)
        reached: 26 of 27 states
```

But, be careful. The result means that all data that is delivered, is delivered in the correct order without deletions etc. We did not check that the data *will* necessarily be delivered. It may be possible for sender and receiver to cycle through a series of states, exchanges erroneous messages, without ever making effective progress. To check this, the state in the sender and in the receiver process that unmistakingly signify progress, were labeled as a ``progress states.'' (In fact, either one by itself would suffice.)

We should now be able to demonstrate the absence of infinite execution cycles that do not pass through any of these progress states. We cannot use the same executable from the last run, but it's easy to setup the verifier for non-progress cycle detection:

```
$ gcc -DNP -o pan pan.c
$ ./pan -l
pan: non-progress cycle (at depth 6)
pan: wrote pan.trail
full statespace search for:
assertion violations and non-progress loops
search was not completed
vector 44 byte, depth reached 8, loops: 1
     12 states, stored
      1 states, linked
      0 states, matched      total: 13
hash conflicts: 0 (resolved)
(size 2^18 states, stack frames: 0/1)
```

There is at least one non-progress cycle. The first one encountered is dumped into the error trail by the verifier, and we can inspect it. The results are shown in the first half of Figure 2. The channel can distort or lose the message infinitely often; true, but not too exciting as an error scenario. To see how many non-progress cycles there are, we can use the `-c` flag. If we set its numeric argument to zero, only a total count of all errors will be printed.

```
$ pan -l -c0
full statespace search for:
assertion violations and non-progress loops
vector 44 byte, depth reached 137, loops: 92
    671 states, stored
      2 states, linked
    521 states, matched      total: 1194
hash conflicts: 39 (resolved)
(size 2^18 states, stack frames: 0/26)
```

There are 92 cases to consider, and we could look at each one, using the `-c` option (`-c1`, `-c2`, `-c3`, ...etc.) But, we can make the job a little easier by at least filtering out the errors caused by infinite message loss. We label all loss events (lines 13, 43, and 48) as progress states, using label names with the common 8-character prefix ``progress,'' and look at the cycles that remain. (Labels go behind the ``::'' flags.)

```
$ spin -a ABP1
$ gcc -DNP -o pan pan.c
```

```
$ ./pan -l
pan: non-progress cycle (at depth 133)
pan: wrote pan.trail

full statespace search for:
assertion violations and non-progress loops
search was not completed

vector 44 byte, depth reached 136, loops: 1

     148 states, stored
       2 states, linked
       2 states, matched       total: 152

hash conflicts: 0 (resolved)
(size 2^18 states, stack frames: 0/26)
```

This time, the trace reveals an honest and a serious bug in the protocol. The second half of Figure 2 shows the trace-back.

```
$ spin -t -r -s ABP0
<<<<>>>>
proc  1 (sender)   line  13, Send err,0,0 -> queue 2 (out)
proc  2 (receiver) line  55, Recv err,0,0 <- queue 2 (in)
proc  2 (receiver) line  56, Send nak,0,0 -> queue 1 (out)
proc  1 (sender)   line  19, Recv nak,0,0 <- queue 1 (in)
spin: trail ends after 12 steps
step 12, #processes: 3
          _p[0] = 6
proc  2 (receiver)      line 36 (state 21)
proc  1 (sender) line 11 (state 6)
proc  0 (_init)  line 67 (state 4)
3 processes created
$
$ spin -t -r -s ABP1
...
proc  2 (receiver) line  39, Recv mesg,0,0 <- queue 2 (in)
proc  2 (receiver) line  47, Send err,0,0  -> queue 1 (out)
proc  1 (sender)   line  20, Recv err,1,0  <- queue 1 (in)
proc  1 (sender)   line  12, Send mesg,0,0 -> queue 2 (out)
proc  2 (receiver) line  39, Recv mesg,0,0 <- queue 2 (in)
proc  2 (receiver) line  52, Send nak,0,0  -> queue 1 (out)
proc  1 (sender)   line  21, Recv nak,0,0  <- queue 1 (in)
proc  1 (sender)   line  12, Send mesg,0,0 -> queue 2 (out)
proc  2 (receiver) line  39, Recv mesg,0,0 <- queue 2 (in)
proc  2 (receiver) line  52, Send nak,0,0  -> queue 1 (out)
<<<<>>>>
proc  1 (sender)   line  21, Recv nak,0,0  <- queue 1 (in)
proc  1 (sender)   line  12, Send mesg,0,0 -> queue 2 (out)
proc  2 (receiver) line  39, Recv mesg,0,0 <- queue 2 (in)
proc  2 (receiver) line  52, Send nak,0,0  -> queue 1 (out)
spin: trail ends after 226 steps
...

    Figure 2.  Error Trails - Extended Alternating Bit Protocol
```

After a single positive acknowledgement is distorted and transformed into an `err` message, sender and receiver get caught in an infinite cycle, where the sender will stubbornly repeat the last message for which it did not receive an acknowledgement, and the receiver, just as stubbornly, will reject that message with a negative acknowledgment.

Digging Deeper

This manual can only give an outline of the main features of *Spin*, and the more common ways in which it can be used for verifications. There is a number of *Spin* features that have not been discussed here, but that may be useful for tackling verification problems.

Spin also allows for a straightforward verification of `tasks,' or `requirements,' modeled as *never*-claims. That is, if the user formalizes a task that is claimed to be performed by the system, *Spin* can quickly either prove or disprove that claim. A never-claim is equivalent to a Büchi Automaton, and can thus model any linear time temporal logic formula. In the newer versions of *Spin*, the user can use normal LTL syntax to define the requirements, and let *Spin* perform the translation chore to the corresponding never-claim.

Spin also allows the user to formalize `reductions' of the system state space, which can be used to restrict a search it to a user defined subset (again, using never claims, this time to `prune' the statespace). With this method it becomes possible to verify quickly whether or not a given error pattern is within the range of behaviors of a system, even when a complete verification is considered to be infeasible.

For details about these alternative uses of `Promela` and the *Spin* software, refer to [5]. The extensions to *Spin* with version 2.0 and 3.0 are outlined in WhatsNew.html.

References

[1] Dijkstra, E.W., ``Guarded commands, non-determinacy and formal derivation of programs.'' *CACM* 18, 8 (1975), 453-457.
[2] Dijkstra, E.W., ``Solution to a problem in concurrent programming control.'' *CACM* 8, 9 (1965), 569.
[3] Hoare, C.A.R., ``Communicating Sequential Processes.'' *CACM* 21, 8 (1978), 666-677.
[4] Holzmann, G.J., ``Algorithms for automated protocol verification.'' *AT&T Technical Journal* 69, 1 (Jan/Feb 1990). Special Issue on Protocol Testing, Specification, Verification.
[5] Holzmann, G.J., *The Spin Model Checker: Primer and Reference Manual* (2003), Addison-Wesley.
[6] Kernighan, B.W. and Ritchie, D.M., *The C Programming Language*. 2nd ed. (1988), Prentice Hall.

Spin Online References
Promela Manual Index
Promela Grammar
Spin HomePage

(Page Updated: 3 June 2007)